Part of speech effect on vision-language representation learning

by

Pasit Tiwawongrut

A Thesis Submitted in Partial Fulfillment of the Requirements for the Degree of
Master of Science in Data Science and Artificial Intelligence

Examination Committee:	Dr. Chaklam Silpasuwanchai
	Dr. Chantri Polprasert
	Dr. Attaphongse Taparugssanagorn

Nationality:	Thai
Previous Degree:	Bachelor of Computer Engineering
	Khon Kaen University
	Thailand

Scholarship Donor:	Asian Institute of Technology

Asian Institute of Technology
School of Engineering and Technology
Thailand
December 2024

ACKNOWLEDGMENTS

ABSTRACT

Vision language (VL) models have shown promising performance across multiple tasks in both zero-shot and fine-tuning setups. Most studies use masked language modeling as a pre-training task, applying random masking to image caption tokens. However, random token masking is not an optimal strategy for training VL models, and effective masking strategies in VL remain underexplored. In this work, we investigate the effects of part of speech (POS) masking, as each POS category contributes differently to sentence meaning. By pre-training models with different POS masking strategies, we evaluate each model on image-text retrieval and visual question answering tasks, categorizing each question type following the VALSE. Our findings contribute to a deeper understanding of how POS masking influences model performance, providing insights that can lead to more effective pre-training strategies for future VL models.

CONTENTS

LIST OF TABLES

LIST OF FIGURES

CHAPTER 1

INTRODUCTION

1.1 Background

Vision language (VL) models have gained significant attention due to their ability to perform both zero-shot and transfer learning, achieving high performance across numerous downstream tasks through pre-training with web-scale image-text pairs (Mo, Kim, Lee, & Shin, 2024; Z. Wang, Wu, Agarwal, & Sun, 2022; J. Zhang, Huang, Jin, & Lu, 2024). Many VL models incorporated masked language modeling (MLM) as a pre-training task, making it an important method to train VL models (J. Li et al., 2021; C. Li et al., 2022; Chen et al., 2020; W. Wang et al., 2023; Tan & Bansal, 2019). Typically, a subset of word tokens is randomly masked at a percentage during training, and the model is tasked with predicting these masked tokens using information from both visual and language modalities. This masking approach has proven to enhance the alignment between visual and linguistic representations, boosting performance in VL tasks (Tan & Bansal, 2019).

Despite the widespread adoption of MLM in VL training, its effects on model performance, efficiency, and training loss remain underexplored. Bitton, Stanovsky, Elhadad, and Schwartz (2021) demonstrated that many of the randomly masked tokens are often stop-words or punctuation, which the model can easily learn without any need for masking. Another study by Wilf et al. (2023) demonstrated that selectively masking infrequent words from the pre-training dataset can boost model performance on out-of-domain datasets during continued training. Additionally, Tou and Sun (2024) suggested that random masking causes the model to rely heavily on local text signals, and it results in inefficient and inconsistent interactions between modalities, leading to suboptimal performance. These find-

ings emphasize the importance of strategic token selection in MLM to enhance VL model performance and efficiency.

In this work, we aimed to address the gap in understanding how masking each part of speech (POS) impacts VL models. Each POS contributes distinctively to sentence meaning: nouns typically denote objects, while verbs describe actions and often demand contextual comprehension. By selectively masking different parts of speech, we could better understand how each POS category affects the alignment between visual and linguistic information. To further explore the effect of each POS, we experimented with fine-tuning the pre-trained model by masking each POS category to assess the effect in the continued training situation. The experiment is designed to answer the following questions:

1. How does masking each POS impact the performance, efficiency, and training loss of VL pre-training models?
2. How does each POS masking strategy affect visual question answering (VQA) performance when analysed based on different question types?
3. What are the effects of part-of-speech masking during the fine-tuning phase of pre-trained VL models?

1.2 Objective

The objectives for our experiment are as listed.

1. Develop a pre-trained VL model to evaluate the impact of masking each POS on performance and training efficiency.
2. Benchmark the performance of each POS masking approach using specialized datasets to gain a deeper understanding of masking effects with retrieval and question answering tasks.
3. Fine-tune the pre-trained VL model to assess the performance and efficiency in the continued training situation.

1.3 Scope

1. The training and testing datasets are web-scale image-text pairs.
2. The model architecture is a cross-attention model, chosen for its ability to jointly predict answers based on information from multiple modalities.

CHAPTER 2

LITERATURE REVIEW

This section of the literature review is organized around two key topics relevant to our study. The first topic addresses VL models, providing an overview of the model architectures recently used in VL models and discussing the choice of the base architecture for the VL model used in this research. The second topic is MLM, an important pre-training approach that has improved VL model performance. Together, these sections provide a comprehensive overview of the methodological foundations of this study.

2.1 Vision-Language Representation Learning

VL learning aims to align visual and linguistic information for multimodal tasks that require reasoning across both modalities, such as image captioning, visual question answering, and multimodal retrieval. The training objective can be divided roughly into three main categories: contrastive, generative, and alignment. Firstly, the contrastive learning objective trains VL representations by maximizing the alignment score between paired images and text while minimizing the score between unpaired images and text (Radford et al., 2021; Jia et al., 2021; J. Yang et al., 2022). Secondly, the generative learning objective focuses on reconstructing masked tokens in either the image or text modality, or both, to learn VL representations (Singh et al., 2022; J. Li et al., 2021; Alayrac et al., 2022). This objective requires the model to utilize both modalities to reconstruct missing tokens, which enhances alignment. Lastly, the alignment objective involves learning VL representations by predicting whether an image and text pair match (Bao et al., 2022). The combined use of these three training objectives has proven effective and is commonly applied across various pre-trained VL models.

Recent advancements in VL fusion methods can be roughly categorized into three main approaches. The first approach is a separate unimodal encoder for each modality, as seen in models like CLIP (Radford et al., 2021) and ALIGN (Jia et al., 2021). This method is trained with the objective of aligning the intermediate outputs of each modality's encoding. The second method uses a cross-attention layer to fuse multimodal inputs, e.g., Flamingo (Alayrac et al., 2022), LXMERT (Tan & Bansal, 2019), and ALBEF (J. Li et al., 2021). The cross-attention layer enables the model to fuse each modality more deeply. Finally, the third approach uses a single large attention model with concatenated image and text tokens as input, as in BEIT-3 (W. Wang et al., 2023), OSCAR (X. Li et al., 2020), UNITER (Chen et al., 2020), FLAVA (Singh et al., 2022), and mPLUG (C. Li et al., 2022). This approach allows for early-stage fusion of each modality, though it requires the highest amount of computational resources. In this work, we adopt the cross-attention method as the base model due to its effectiveness in fusing multimodal inputs. Additionally, this approach allows the model to be trained using the MLM task. We also use all three training objectives with a modified MLM for this experiment.

2.2 Masked Language Modelling

MLM is a widely used pre-training method in language model (LM) training (Devlin, Chang, Lee, & Toutanova, 2018; Lan, 2019; Yu et al., 2022; S. Zhang et al., 2022; Guu, Lee, Tung, Pasupat, & Chang, 2020) as a self-supervised task. BERT (Devlin et al., 2018) proposed MLM as a pre-training task, which has been proven effective for pre-training language models. The MLM task involves replacing some input tokens with a special [MASK] token, and the model must predict the masked tokens based on the given unmasked tokens. In the field of VL models, many VL models have also adopted MLM as a training task to train the model to predict masked text based on visual information (J. Li et al., 2021; C. Li

et al., 2022; Chen et al., 2020; W. Wang et al., 2023).

In the field of selective masking strategies in natural language processing, several works have further refined MLM to enhance training efficiency. ERNIE (Sun et al., 2019), SpanBERT (Joshi et al., 2020), and n-gram Masking (Levine et al., 2021) propose span masking instead of single-token masking, which forces the model to rely more on long-range dependencies rather than adjacent tokens, resulting in better performance compared to BERT (Devlin et al., 2018). Considering linguistic features, D. Yang, Zhang, and Zhao (2023) conducted a training analysis based on POS masking focused on LM training. The results showed that focusing the masking of non-function words (ADJ, ADV, NOUN, PROPN, and VERB) in the later stages of training can encourage the LM model to develop a better contextual understanding.

For selective masking in VL training, Bitton et al. (2021) introduced an object token masking strategy, selectively masking object tokens in image captions and pre-training the model. This approach achieved superior performance compared to random masking. Another study by Wilf et al. (2023) showed that selectively masking infrequent words from the pre-training dataset during continued training enhances model performance on out-of-domain datasets. Additionally, (Tou & Sun, 2024) proposed a curriculum-based masking strategy in which a reinforcement learning agent dynamically selects masking spans based on cross-modal interactions. This method improved the model's multimodalities understanding while reducing the dataset size needed for effective training. In this work, we conduct experiments to analyze the impact of each POS on results within a VL setting.

CHAPTER 3

METHODOLOGY

In this chapter, the methodology is detailed as follows. First, we describe the architecture of the model. Second, we explain all pre-training loss functions used in this experiment. Third, the details of POS tagging are provided. Fourth, we outline the datasets used in this experiment. Lastly, we provide details on the visual question answering setup.

Figure 3.1
Overall methodology

Pre-training the model with a MLM task by masking tokens based on the POS in the image captions.

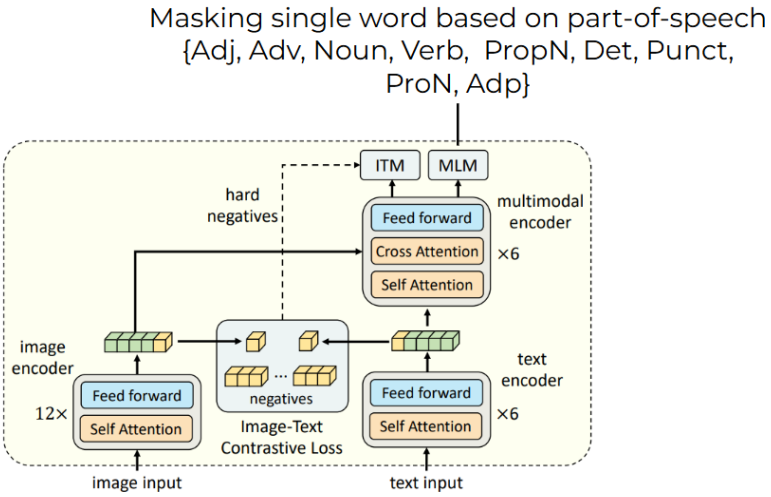

3.1 Model architecture

As shown in Figure 3.1, our model includes three main components: an image encoder, a text encoder, and a multimodal encoder. The first component is the

image encoder, for which we use ViT (Dosovitskiy et al., 2021), modified following (Radford et al., 2021), as the image encoder in this experiment. The second component is the text encoder, which employs a transformer architecture as BERT (Devlin et al., 2018) to encode image captions with BERT tokenizer for tokenization. The final component is the multimodal encoder, where VL interactions occur.

Given a training dataset D consisting of image-text pairs $(I_i, T_i) \in D$, where I_i is the image and T_i is the image caption of the i-th image, each image is first encoded as a sequence of tokens $\{v_{cls}, v_1, \ldots, v_n\}$ using ViT (Dosovitskiy et al., 2021). Here, v_{cls} represents the embedding of the [CLS] token prepended to the image patch sequence. In this experiment, the image encoder was initialized with ViT-B-32 pre-trained on ImageNet-21K (Deng et al., 2009). Next, we use a 6-layer transformer, randomly initialized, to encode the image caption T_i into text embeddings $\{w_{cls}, w_1, \ldots, w_n\}$, where w_{cls} is the embedding of the [CLS] token. Finally, both text and image encodings are passed through the multimodal encoder to fuse both inputs, producing multimodal encodings. For the multimodal encoder, a cross-attention layer is used, where both keys and values are the image encodings, and the text encoding serves as the query in the cross-attention layer.

3.2 Pre-training Objectives

In this work, we pre-train our model with three objectives: masked language modeling (MLM), image-text constrative learning (ITC) and image-text matching (ITM).

3.2.1 Mask Language Modelling

Typically, a percentage of tokens $\{w_1, \ldots, w_T\}$ are replaced with a special [MASK] token to create a masked caption T^{mask}. However, in this work, the masked tokens were selected based on POS type instead of randomly masking. The model

trained to predict the original tokens at the masked positions, conditioned on both the unmasked tokens in T^{mask} and the visual features of I as $p^{\text{mask}}(I, T^{\text{mask}})$. Let y^{mask} be a one-hot vector representing the ground-truth vocabulary for the masked token, where the masked token has a probability of 1. The model's objective is to minimize the cross-entropy $\mathbf{H}$, given by:

$$\mathcal{L}_{\text{MLM}} = \mathbf{H}(y^{\text{mask}}, p^{\text{mask}}(I, T^{\text{mask}})))$$

3.2.2 Image-Text Contrastive Learning

To improve each unimodal encoder's representation, we used image-text constrative learning to improve alignment of each modality. ITC aims to improve alignment by maximizing the similarity score of image and text from the same pair with the score function $s(I,T) = v_{cls}^{\top} w_{cls}$, and minimizing the similarity score of image and text not from its pair. We then calculate the softmax-normalized similarity score for each image to any text and each text to any image, identified as image-to-text $p^{i2t} \in \mathbb{R}^M$ and text-to-image $p^{t2i} \in \mathbb{R}^M$ scores as:

$$p_i^{i2t}(I) = \frac{\exp\left(s(I,T_i)\right)/\tau}{\sum_{m=1}^{M} \exp\left(s(I,T_m)\right)/\tau}, \quad p_i^{t2i}(T) = \frac{\exp\left(s(T,I_i)\right)/\tau}{\sum_{m=1}^{M} \exp\left(s(T,I_m)\right)/\tau}$$

where τ is a learnable temperature parameter. Let $y^{i2t}(I) \in \{0,1\}^M$ and $y^{t2i}(T) \in \{0,1\}^M$ be a ground truth with probability of 1 at a position of the same pair, and probability of 0 on the other hand. The ITC loss is calculated as cross-entropy $\mathbf{H}$ between p and y:

$$\mathcal{L}_{\text{ITC}} = \frac{1}{2}(\mathbf{H}(y^{i2t}, p^{i2t}) + \mathbf{H}(y^{t2i}, p^{t2i}))$$

3.2.3 Image-Text Matching

To further improve multimodal alignment in the VL model, image-text matching was employed to enhance alignment. The model is trained to predict whether an image and caption are from the same pair. A fully connected layer, followed by

a softmax function, is added over the model. This layer takes the [CLS] embedding from the multimodal encoding as input to predict whether the pair is positive (matched) or negative (unmatched).

The loss function for ITM, using cross-entropy loss, is defined as:

$$\mathcal{L}_{\text{ITM}} = \mathbf{H}(y^{\text{itm}}, p^{\text{itm}}(I,T)),$$

where y^{itm} is a one-hot ground-truth label, and $p^{\text{itm}}(I,T)$ is the predicted class probability.

The full pre-training objective of our work can be written as:

$$\mathcal{L} = \mathcal{L}_{\text{MLM}} + \mathcal{L}_{\text{ITC}} + \mathcal{L}_{\text{ITM}}$$

3.3 Part Of Speech Masking

In this experiment, we explored the effect of each POS on VL learning in terms of performance, efficiency, and training loss. For each image caption, each token was classified into POS categories for masking. We used POS-tagging tools SpaCy[1] to classify each word into POS classes based on the Universal POS tag set[2].

3.4 Pre-Training Dataset

We pre-trained the model on the Conceptual Captions dataset (Sharma, Ding, Goodman, & Soricut, 2018), which consists of 3.3 million image-text pairs. In Conceptual Captions dataset, an automated process was used to select, filter, and refine these image-caption pairs to ensure they are clear, informative, and suitable for effective model training.

[1]POS-tagging tool SpaCy: https://spacy.io/

[2]Universal POS tag set: https://universaldependencies.org/u/pos/

3.5 Continue Training

In this experiment, we explored the effect of each POS in the continue training situation, where the amount of dataset is limited, and the domain is difference from the pre-training dataset. We trained each model on difference categories of POS with every datasets, including GoodNews (Biten, Gomez, Rusinol, & Karatzas, 2019), RSICD (Lu, Wang, Zheng, & Li, 2017), and Sketchy Scene (Lian, Yangdong, & Yuejie, 2021) The model was initialized with ALBEF pre-training weights. The GoodNews dataset is an image-text pair dataset gather from New York Times. This dataset have total image-text pairs of 466,000 pairs, randomly splited into 424,000 for training, 18,000 for validation, and 23,000 for testing. For RSICD, the dataset include a remote sensing image in total of 10,921 images with 5 captions per image. The Sketchy Scene dataset include total of 1,000 image-text pairs.

3.6 Evaluation

In this work, we evaluated each model trained with different types of POS masking through image-text retrieval and visual question answering tasks. Details of the evaluation methods and datasets used in these tasks are provided in this section.

3.6.1 Image-Text Retrieval

For the image-text retrieval task, we evaluated the effect of masking on each POS category in two situations: pre-training, and continue training. First, pre-training evaluation, the model was tested by performing zero-shot evaluations on the Flickr30K (Plummer et al., 2015) dataset for both image retrieval (IR) and text retrieval (TR). The Flickr30K dataset is used to assess the model's overall performance in retrieval tasks. The baseline for this evaluation was the model pre-trained with the randomly masking strategy. Second, continue training evalation, we benchmarked each model with GoodNews, RSICD, and Sketchy Scene

dataset. This setup allowed us to analyze how different POS masking strategies affect the model's retrieval performance and the alignment between visual and textual representations.

3.6.2 Image-Text Matching

As demonstrated by Tou and Sun (2024), the results suggest that masking strategies impact a model's ability to understand attributes, relationships, and word order. In this work, we benchmarked each pre-trained model with specific POS masking against two tasks: Attribution, Relation and Order benchmark (Yuksekgonul, Bianchi, Kalluri, Jurafsky, & Zou, 2023) and VALSE benchmark (Parcalabescu et al., 2022). For Attribution, Relation and Order benchmark, this benchmark contains 4 sub-tasks, built over well-known datasets, including Visual Genome Relation (VGR), Visual Genome Attribution (VGA), COCO order (Co), and Filckr order (Fo). This benchmark is designed to assess the VL model's understanding of compositional relationships by swapping and replacing words in image captions, such as altering "The horse is eating the grass" to "The grass is eating the horse." For the VALSE dataset, this benchmark categorizes each image-text sample into different linguistic phenomena, including six distinct types: existence, plurality, counting, relation, action, and coreference. Each image caption in the VALSE dataset also includes a "foil" version, where words related to each caption category are modified. This task is a classification task, where the model has to predict the correct caption for each image. Evaluating models against this benchmark provides valuable insights into their semantic and contextual understanding of vision and language modality.

3.6.3 Visual Question Answering

The visual question answering (VQA) task requires the model to generate answers, which require an additional decoder over the multimodal encoder, as shown in Figure 3.2. In this work, we appended a 6-layer transformer as a decoder, initialized with the parameters of the multimodal encoder. Answers are generated

Figure 3.2
Visual question answering model architecture

Modified model architecture for VQA task.

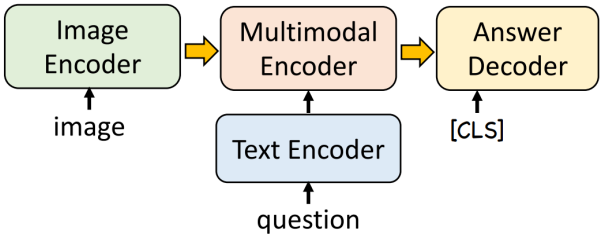

in an auto-regressive manner, with multimodal embeddings as input and a special start-of-sequence token ([CLS]) as the initial input to the decoder. The benchmark dataset for the VQA task is the VQA2.0 dataset (Goyal, Khot, Summers-Stay, Batra, & Parikh, 2017), which is constructed using images from COCO (Lin et al., 2014). This dataset includes 83,000 images for training, 41,000 for validation, and 81,000 for testing. We further train our model using the VQA2.0 training set. We categorized each question into six categories following VALSE linguistic phenomena categories to provide deeper insights.

CHAPTER 4

Results

In this chapter, we provide all the experiment results for each experiment and evaluation.

4.1 Pre-training

4.1.1 Flickr30K

4.2 Image-Text Matching

4.2.1 VALSE

4.3 Visual question answering

4.4 Fine tuning

[1] Spacial Relation

Table 4.1
Flickr30K benchmark image retrieval result.

Masking Method		Flickr30K					
		TR			IR		
		r@1	r@5	r@10	r@1	r@5	r@10
Random Masking		67.00	88.00	93.75	52.61	80.14	87.76
Non-function	Noun	67.15	88.60	94.65	52.73	80.45	87.79
	Verb	54.85	82.85	90.05	43.82	73.84	82.82
	Adj	62.30	87.30	92.40	47.39	75.47	84.06
	Adv	46.85	76.25	85.75	36.40	66.38	76.78
	PropN	44.85	74.40	84.10	34.91	64.09	75.01
Function	Det	71.05	92.00	95.30	56.01	81.93	88.59
	Aux	52.10	79.60	88.20	41.13	70.92	80.68
	ProN	51.45	78.80	87.10	39.97	69.58	79.32
	Adp	65.05	88.25	93.40	51.19	78.83	85.15

Table 4.2
VALSE benchmark for image-text matching result.

Masking Method		VALSE											
		Existence	Prularity	Counting			Sp.Re [1]	Action		Coreference		Foil-it!	Avg
		Quantifiers	Number	Balanced	Small number	Adversarial	Relations	Replacement	Actant swap	Standard	Clean		
Random Masking													
Non-function	Noun												
	Verb												
	Adj												
	Adv												
	PropN												
Function	Det												
	Aux												
	Punct												
	ProN												
	Adp												

Table 4.3
Visual question answering result with VQA2.0 benchmark.

Masking Method		VQA2.0					
		Existence	Prularity	Counting	Sp.Re [2]	Action	Coreference
Random Masking							
Non-function	Noun						
	Verb						
	Adj						
	Adv						
	PropN						
Function	Det						
	Aux						
	Punct						
	ProN						
	Adp						

Table 4.4
Fine tuning benchmark.

Masking Method		GoodNews		RSICD		Sketchy Scene	
		R@5					
		TR	IR	TR	IR	TR	IR
Random Masking							
Non-function	Noun						
	Verb						
	Adj						
	Adv						
	PropN						
Function	Det						
	Aux						
	Punct						
	ProN						
	Adp						

CHAPTER 5

DISCUSSION

CHAPTER 6

CONCLUSION

REFERENCES

Alayrac, J.-B., Donahue, J., Luc, P., Miech, A., Barr, I., Hasson, Y., ... others (2022). Flamingo: a visual language model for few-shot learning. *Advances in Neural Information Processing Systems*, *35*, 23716–23736.

Bao, H., Wang, W., Dong, L., Liu, Q., Mohammed, O. K., Aggarwal, K., ... Wei, F. (2022). Vlmo: Unified vision-language pre-training with mixture-of-modality-experts. *Advances in Neural Information Processing Systems*, *35*, 32897–32912.

Biten, A. F., Gomez, L., Rusinol, M., & Karatzas, D. (2019). Good news, everyone! context driven entity-aware captioning for news images. In *Proceedings of the ieee/cvf conference on computer vision and pattern recognition* (pp. 12466–12475).

Bitton, Y., Stanovsky, G., Elhadad, M., & Schwartz, R. (2021). Data efficient masked language modeling for vision and language. *arXiv preprint arXiv:2109.02040*.

Chen, Y.-C., Li, L., Yu, L., El Kholy, A., Ahmed, F., Gan, Z., ... Liu, J. (2020). Uniter: Universal image-text representation learning. In *Computer vision – eccv 2020: 16th european conference, glasgow, uk, august 23–28, 2020, proceedings, part xxx* (p. 104–120). Berlin, Heidelberg: Springer-Verlag. Retrieved from https://doi.org/10.1007/978-3-030-58577-8_7 doi: 10.1007/978-3-030-58577-8_7

Deng, J., Dong, W., Socher, R., Li, L.-J., Li, K., & Fei-Fei, L. (2009). Imagenet: A large-scale hierarchical image database. In *2009 ieee conference on computer vision and pattern recognition* (pp. 248–255).

Devlin, J., Chang, M.-W., Lee, K., & Toutanova, K. (2018). Bert: Pre-training of deep bidirectional transformers for language understanding. *arXiv preprint arXiv:1810.04805*.

Dosovitskiy, A., Beyer, L., Kolesnikov, A., Weissenborn, D., Zhai, X., Unterthiner, T., ... Houlsby, N. (2021). An image is worth 16x16 words: Transformers for image recognition at scale. In *International conference on learning representations*. Retrieved from https://openreview.net/forum?id=YicbFdNTTy

Goyal, Y., Khot, T., Summers-Stay, D., Batra, D., & Parikh, D. (2017). Making the v in vqa matter: Elevating the role of image understanding in visual question answering. In *Proceedings of the ieee conference on computer vision and pattern recognition* (pp. 6904–6913).

Guu, K., Lee, K., Tung, Z., Pasupat, P., & Chang, M.-W. (2020). *Realm: Retrieval-augmented language model pre-training.* Retrieved from https://arxiv.org/abs/2002.08909

Jia, C., Yang, Y., Xia, Y., Chen, Y.-T., Parekh, Z., Pham, H., ... Duerig, T. (2021). Scaling up visual and vision-language representation learning with noisy text supervision. In *International conference on machine learning* (pp. 4904–4916).

Joshi, M., Chen, D., Liu, Y., Weld, D. S., Zettlemoyer, L., & Levy, O. (2020). SpanBERT: Improving pre-training by representing and predicting spans. *Transactions of the Association for Computational Linguistics*, *8*, 64–77. Retrieved from https://aclanthology.org/2020.tacl-1.5 doi: 10.1162/tacl_a_00300

Lan, Z. (2019). Albert: A lite bert for self-supervised learning of language representations. *arXiv preprint arXiv:1909.11942*.

Levine, Y., Lenz, B., Lieber, O., Abend, O., Leyton-Brown, K., Tennenholtz, M., & Shoham, Y. (2021). {PMI}-masking: Principled masking of correlated spans. In *International conference on learning representations.* Retrieved from https://openreview.net/forum?id=3Aoft6NWFej

Li, C., Xu, H., Tian, J., Wang, W., Yan, M., Bi, B., ... others (2022). mplug: Effective and efficient vision-language learning by cross-modal skip-connections. *arXiv preprint arXiv:2205.12005*.

Li, J., Selvaraju, R., Gotmare, A., Joty, S., Xiong, C., & Hoi, S. C. H. (2021). Align before fuse: Vision and language representation learning with momentum distillation. *Advances in neural information processing systems*, *34*, 9694–9705.

Li, X., Yin, X., Li, C., Zhang, P., Hu, X., Zhang, L., ... others (2020). Oscar: Object-semantics aligned pre-training for vision-language tasks. In *Computer vision–eccv 2020: 16th european conference, glasgow, uk, august 23–28, 2020, proceedings, part xxx 16* (pp. 121–137).

Lian, Z., Yangdong, C., & Yuejie, Z. (2021, August). Sketchy scene captioning: Learning multi-level semantic information from sparse visual scene cues. In S. Li et al. (Eds.), *Proceedings of the 20th chinese national conference on computational linguistics* (pp. 1167–1177). Huhhot, China: Chinese Information Processing Society of China. Retrieved from https://aclanthology.org/2021.ccl-1.104

Lin, T.-Y., Maire, M., Belongie, S., Hays, J., Perona, P., Ramanan, D., ... Zitnick, C. L. (2014). Microsoft coco: Common objects in context. In *Computer vision–eccv 2014: 13th european conference, zurich, switzerland, septem-*

ber 6-12, 2014, proceedings, part v 13 (pp. 740–755).

Lu, X., Wang, B., Zheng, X., & Li, X. (2017). Exploring models and data for remote sensing image caption generation. *IEEE Transactions on Geoscience and Remote Sensing*, *56*(4), 2183–2195.

Mo, S., Kim, M., Lee, K., & Shin, J. (2024). S-clip: Semi-supervised vision-language learning using few specialist captions. *Advances in Neural Information Processing Systems*, *36*.

Parcalabescu, L., Cafagna, M., Muradjan, L., Frank, A., Calixto, I., & Gatt, A. (2022, May). VALSE: A task-independent benchmark for vision and language models centered on linguistic phenomena. In *Proceedings of the 60th annual meeting of the association for computational linguistics (volume 1: Long papers)* (pp. 8253–8280). Dublin, Ireland: Association for Computational Linguistics. Retrieved from `https://aclanthology.org/2022.acl-long.567`

Plummer, B. A., Wang, L., Cervantes, C. M., Caicedo, J. C., Hockenmaier, J., & Lazebnik, S. (2015). Flickr30k entities: Collecting region-to-phrase correspondences for richer image-to-sentence models. In *Proceedings of the ieee international conference on computer vision* (pp. 2641–2649).

Radford, A., Kim, J. W., Hallacy, C., Ramesh, A., Goh, G., Agarwal, S., ... others (2021). Learning transferable visual models from natural language supervision. In *International conference on machine learning* (pp. 8748–8763).

Sharma, P., Ding, N., Goodman, S., & Soricut, R. (2018). Conceptual captions: A cleaned, hypernymed, image alt-text dataset for automatic image captioning. In *Proceedings of acl.*

Singh, A., Hu, R., Goswami, V., Couairon, G., Galuba, W., Rohrbach, M., & Kiela, D. (2022). Flava: A foundational language and vision alignment model. In *Proceedings of the ieee/cvf conference on computer vision and pattern recognition* (pp. 15638–15650).

Sun, Y., Wang, S., Li, Y., Feng, S., Chen, X., Zhang, H., ... Wu, H. (2019). ERNIE: enhanced representation through knowledge integration. *CoRR*, *abs/1904.09223*.

Tan, H., & Bansal, M. (2019). Lxmert: Learning cross-modality encoder representations from transformers. In *Proceedings of the 2019 conference on empirical methods in natural language processing.*

Tou, K., & Sun, Z. (2024, June). Curriculum masking in vision-language pretraining to maximize cross modal interaction. In K. Duh, H. Gomez, & S. Bethard (Eds.), *Proceedings of the 2024 conference of the north amer-*

ican chapter of the association for computational linguistics: Human language technologies (volume 1: Long papers) (pp. 3672–3688). Mexico City, Mexico: Association for Computational Linguistics. Retrieved from https://aclanthology.org/2024.naacl-long.203 doi: 10.18653/v1/2024.naacl-long.203

Wang, W., Bao, H., Dong, L., Bjorck, J., Peng, Z., Liu, Q., ... others (2023). Image as a foreign language: Beit pretraining for vision and vision-language tasks. In *Proceedings of the ieee/cvf conference on computer vision and pattern recognition* (pp. 19175–19186).

Wang, Z., Wu, Z., Agarwal, D., & Sun, J. (2022). *Medclip: Contrastive learning from unpaired medical images and text.* Retrieved from https://arxiv.org/abs/2210.10163

Wilf, A., Akter, S. N., Mathur, L., Liang, P. P., Mathew, S., Shou, M., ... Morency, L.-P. (2023). Difference-masking: Choosing what to mask in continued pretraining. *arXiv preprint arXiv:2305.14577*.

Yang, D., Zhang, Z., & Zhao, H. (2023). *Learning better masking for better language model pre-training.*

Yang, J., Li, C., Zhang, P., Xiao, B., Liu, C., Yuan, L., & Gao, J. (2022). Unified contrastive learning in image-text-label space. In *Proceedings of the ieee/cvf conference on computer vision and pattern recognition* (pp. 19163–19173).

Yu, W., Zhu, C., Fang, Y., Yu, D., Wang, S., Xu, Y., ... Jiang, M. (2022). *Dict-bert: Enhancing language model pre-training with dictionary.* Retrieved from https://arxiv.org/abs/2110.06490

Yuksekgonul, M., Bianchi, F., Kalluri, P., Jurafsky, D., & Zou, J. (2023). When and why vision-language models behave like bags-of-words, and what to do about it? In *International conference on learning representations.* Retrieved from https://openreview.net/forum?id=KRLUvxh8uaX

Zhang, J., Huang, J., Jin, S., & Lu, S. (2024). Vision-language models for vision tasks: A survey. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, *46*(8), 5625-5644. doi: 10.1109/TPAMI.2024.3369699

Zhang, S., Roller, S., Goyal, N., Artetxe, M., Chen, M., Chen, S., ... Zettlemoyer, L. (2022). *Opt: Open pre-trained transformer language models.* Retrieved from https://arxiv.org/abs/2205.01068